

REGRESSION ANALYSIS-2 PROJECT REPORT

DATA ANALYSIS USING PISA
STUDENT DATA

Lecture: Regression Analysis 2
Teacher: Gülhayat Gölbaşı Şimşek

Metin Mirza VATANSEVER - 20023055
Münir Göktan SOYKURT - 20023048
Enes COŞKUN - 21023005

ABOUT PROJECT

The main goal of this project is to analyse PISA math score of Great Britian students. Analyse will be done by using various regression models in R studio.

ABOUT DATA

The given data is considered a big one. The data contains too many variables. The first step of this process is to analyse the data and choose the variables to work on. We will walk through the variables will be needed during this analysis. To work on this complicated data we create new variables with given variables.

```
# RSN = (PV1MPRE+...+PV10MPRE)/10 - Score of math reasoning
# UNDR = (PV1MPIN+...+PV10MPIN)/10 - Score of interpreting applying and evaluating of math
# MUNDR = (PV1MPFS+...+PV10MPFS)/10 - Score of math formulation
# EMPL = (PV1MPEM+...+PV10MPEM)/10 - Score of employing math concepts
# CRT = (PV1MCUD+...+PV10MCUD)/10 - Score of uncertainty and data
# SPC = (PV1MCSS+...+PV10MCSS)/10 - Score of space and shape understanding
# MATH = (PV1MATH+...+PV10MATH)/10 - Score of math
```

SUMMARY OF THE DATA

```
> describe(dataset)
```

	vars	n	mean	sd	median	trimmed	mad	min	max	range	skew
CNT	1	12972	1.00	0.00	1	1.00	0.00	1	1	0	NaN
CNTRYID	2	12972	826.00	0.00	826	826.00	0.00	826	826	0	NaN
CNTSCHID	3	12972	82612714.48	21624.71	82600257	82609628.39	192.74	82600001	82650120	50119	1.15
CNTSTUID	4	12972	82618948.52	19761.23	82610342	82616831.94	7652.44	82600003	82654772	54769	1.02
CYC	5	12972	1.00	0.00	1	1.00	0.00	1	1	0	NaN
NatCen	6	12972	82725.11	43.37	82700	82718.89	0.00	82700	82800	100	1.15
STRATUM	7	12972	17.95	12.96	13	17.41	14.83	1	39	38	0.35
SUBNATIO	8	12972	8260502.16	867.30	8260000	8260377.72	0.00	8260000	8262000	2000	1.15
REGION	9	12972	82613.84	3.64	82612	82613.42	1.48	82611	82620	9	1.02
OECD	10	12972	1.00	0.00	1	1.00	0.00	1	1	0	NaN
ADMINMODE	11	12972	2.00	0.00	2	2.00	0.00	2	2	0	NaN
LANGTEST_QQQ	12	11537	314.98	11.26	313	313.00	0.00	313	379	66	5.51
LANGTEST_COG	13	12972	319.61	57.08	313	313.00	0.00	313	979	666	11.08
ROOKTN	14	12972	13.77	7.71	13	13.00	8.90	1	36	35	0.78

The data contains a lot of variables. The output is too big to fit in so we put a little piece of the output.

```
describe(dataset$MATH)
```

vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
1	12972	481.82	91.52	480.74	480.9	95.88	182.77	834.49	651.72	0.1	-0.34	0.8

```
describe(dataset$SCIE)
```

vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
1	12972	492.74	96.77	492.23	492.21	103.25	193.37	837.17	643.8	0.06	-0.43	0.85

```
describe(dataset$READ)
```

vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
1	12972	490.45	97.04	492.52	491.7	100.79	154.41	816.35	661.95	-0.1	-0.28	0.85

```
describe(dataset$CRT)
```

vars	n	mean	sd	median	trimmed	mad	min	max	range	skew	kurtosis	se
1	12972	490.34	99.85	488.78	489.47	103.47	168.18	889.8	721.62	0.09	-0.27	0.88

Output of descriptive statistics of variables that we will be use in regression models.

04

```
summary(dataset$MATH)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
182.8	415.9	480.7	481.8	545.1	834.5

```
summary(dataset$SCIE)
```


Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
193.4	422.6	492.2	492.7	561.8	837.2

```
summary(dataset$READ)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
154.4	423.9	492.5	490.4	559.8	816.4

Summary of the dependent and independent variables that we will use in regression models.

```
par(mfrow=c(1,3))  
boxplot(dataset$MATH, col = "lightblue", main = "MATH")  
boxplot(dataset$SCIE, col = "lightyellow", main = "SCIE")  
boxplot(dataset$READ, col = "lightgreen", main = "READ")  
par(mfrow=c(1,1))
```



The box plot of the variables. So we can observe the outliers

This chapter is about preperation for convinient use of data in regression models. The given rule was to choose random two students from each school.

```
schoolid <- str(dataset$CNTSCHID)
schoolid <- as.character(dataset$CNTSCHID)
schoolid
dataset_nondup <- dataset[!duplicated(dataset$CNTSCHID), ]
school_counts <- dataset_nondup %>%
  count(CNTSCHID, sort = TRUE)
print(school_counts)
```

This code gives us how many unique school ids PISA data has.

```
set.seed(123)
random_students <- dataset %>%
  group_by(CNTSCHID) %>%
  sample_n(size = 2, replace = FALSE) %>%
  ungroup()
```


This code randomly chooses two students from each unique school id. So with this code we create our data that we will work on.

```
> random_students
# A tibble: 902 x 657
  CNT CENTRYID CNTSCHID NTSTUID CYC NatCen STRATUM SUBNATIO REGION OECD ADMINMODE LANGTEST_QQQ LANGTEST_COG
  <chr> <dbl> <dbl> <dbl> <chr> <dbl> <chr> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 GBR 826 82600001 2609582 08MS 82700 QUK02 8260000 82611 1 2 313 313
2 GBR 826 82600001 2612396 08MS 82700 QUK02 8260000 82611 1 2 313 313
3 GBR 826 82600002 2604924 08MS 82700 QUK03 8260000 82611 1 2 313 313
4 GBR 826 82600002 2601162 08MS 82700 QUK03 8260000 82611 1 2 313 313
5 GBR 826 82600003 2604422 08MS 82700 QUK03 8260000 82611 1 2 313 313
6 GBR 826 82600003 2610261 08MS 82700 QUK03 8260000 82611 1 2 313 313
7 GBR 826 82600004 2609871 08MS 82700 QUK25 8260000 82612 1 2 313 313
8 GBR 826 82600004 2606038 08MS 82700 QUK25 8260000 82612 1 2 313 313
9 GBR 826 82600005 2601525 08MS 82700 QUK02 8260000 82611 1 2 313 313
10 GBR 826 82600005 2612599 08MS 82700 QUK02 8260000 82611 1 2 313 313
# i 892 more rows
# i 644 more variables: BOOKID <dbl>, ST001D01T <dbl>, ST003D02T <dbl>, ST003D03T <dbl>, ST250D06JA <dbl>,
# ST250D07JA <dbl>, ST251Q01JA <dbl>, ST251Q02JA <dbl>, ST251Q03JA <dbl>, ST251Q04JA <dbl>, ST251Q06JA <dbl>,
# ST251Q07JA <dbl>, ST251D08JA <dbl>, ST251D09JA <dbl>, ST253Q01JA <dbl>, ST254Q01JA <dbl>, ST254Q02JA <dbl>,
# ST254Q03JA <dbl>, ST254Q04JA <dbl>, ST254Q05JA <dbl>, ST254Q06JA <dbl>, ST255Q01JA <dbl>, ST256Q01JA <dbl>,
# ST256Q02JA <dbl>, ST256Q03JA <dbl>, ST256Q06JA <dbl>, ST256Q07JA <dbl>, ST256Q08JA <dbl>, ST256Q09JA <dbl>,
# ST256Q10JA <dbl>, ST230Q01JA <dbl>, ST005Q01JA <dbl>, ST006Q01JA <dbl>, ST006Q02JA <dbl>, ...
# i Use `print(n = ...)` to see more rows, and `colnames()` to see all variable names
```

```
dataset$AGE[dataset$AGE == 0] <- NA
dataset$ESCS[dataset$ESCS == 0] <- NA
dataset$HISEI[dataset$HISEI == 0] <- NA
dataset$STRATUM[dataset$STRATUM == 0] <- NA
dataset$PAREDINT[dataset$PAREDINT == 0] <- NA
```

This code fills missing value with 'NA'. So the missing value does not effect the result.

```
correlation_matrix <- cor(dataset[,c("MATH", "READ", "SCIE", "CRT",
                                     "SPC", "EMPL", "ST268Q04JA", "ST059Q01TA")], use = "complete.obs")
corrplot(correlation_matrix, method = "circle", type = "upper",
         title = "Corelation Matrix between math and science",
         cl.pos = "b", tl.col = "black")
```



This output shows relationship between different variables.

07

LINEAR REGRESSION MODEL

Linear Regression is a statistical method used to model the relationship between a dependent variable and one or more independent variables. The method assumes a linear relationship, expressed by the equation $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_n X_n + \epsilon$, where:

- Y is the dependent variable you're predicting.
- $(\beta_1, \beta_2, \dots, \beta_n)$ are coefficients indicating the importance of each independent variable.
- (X_1, X_2, \dots, X_n) are the independent variables.
- (ϵ) is the error term, representing what the model can't explain.

Uses of Linear Regression:

Prediction and Forecasting: Useful in predicting outcomes based on changes in predictor variables.

Data Inference: Helps understand which factors influence the dependent variable and how strongly they do so.

Simplicity and Interpretability: The model's straightforward nature allows easy interpretation of the results.

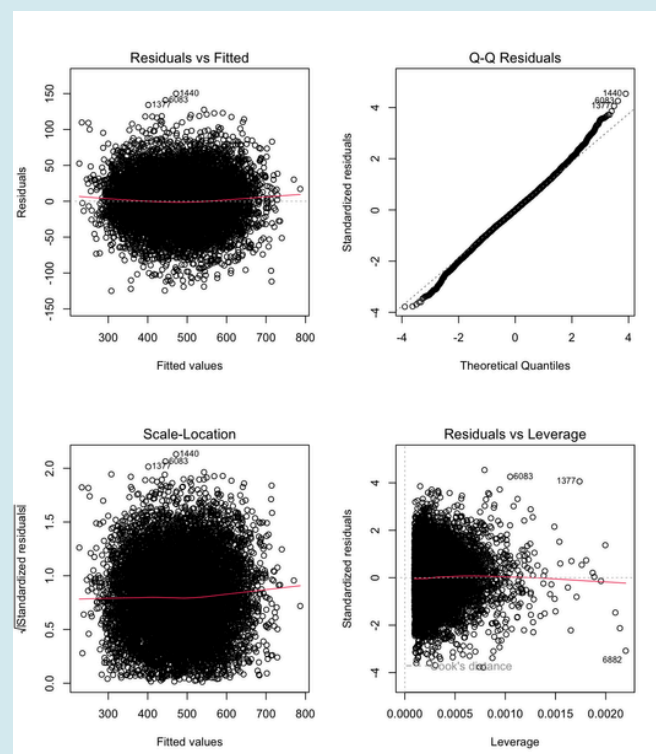
Foundation for Other Methods: Linear regression underpins more complex analytical methods and machine learning algorithms.

Why Use It?

It quantifies the strength of relationships between the dependent and independent variables.

It offers clarity and simplicity in understanding data relationships.

It's instrumental in decision-making processes where predictors need quantification.



```
set.seed(12345)
train_index <- createDataPartition(dataset$MATH, p = 0.8, list = FALSE)
train_set <- dataset[train_index, ]
test_set <- dataset[-train_index, ]
```

```
model <- lm(MATH ~ READ + SCIE, data = train_set)
train_predictions <- predict(model, train_set)
test_predictions <- predict(model, test_set)
train_rmse <- sqrt(mean((train_set$MATH - train_predictions)^2))
test_rmse <- sqrt(mean((test_set$MATH - test_predictions)^2))
cat("Train RMSE:", train_rmse, "\n")
cat("Test RMSE:", test_rmse, "\n")
summary(model)
par(mfrow = c(2, 2))
plot(model)
```

```
> summary(model)
```

Train RMSE: 26.7187
Test RMSE: 33.42844

Call:
lm(formula = MATH ~ READ + SCIE, data = train_set)

Residuals:

Min	1Q	Median	3Q	Max
-124.987	-20.957	-0.244	20.929	150.084

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	36.624147	1.726725	21.21	<2e-16 ***
READ	0.249706	0.006870	36.35	<2e-16 ***
SCIE	0.654898	0.006887	95.10	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 33.07 on 10377 degrees of freedom
Multiple R-squared: 0.8702, Adjusted R-squared: 0.8702
F-statistic: 3.479e+04 on 2 and 10377 DF, p-value: < 2.2e-16

08 LINEAR MODEL GRAPHS EXPLANATION

1. Residuals vs Fitted

Description: This plot shows residuals (differences between observed and predicted values) against fitted values. Ideally, residuals should be randomly dispersed around the central line.

Observation: Residuals appear to show a pattern around the central line, suggesting potential non-linearity or failure to capture all variability in the data.

2. Q-Q Residuals

Description: This plot checks if standardized residuals follow a normal distribution. Points should mostly lie on the reference line.

Observation: Most residuals align well with the theoretical line, indicating normal distribution, but outliers are evident at the extremes, suggesting potential issues with extreme values.

3. Scale-Location

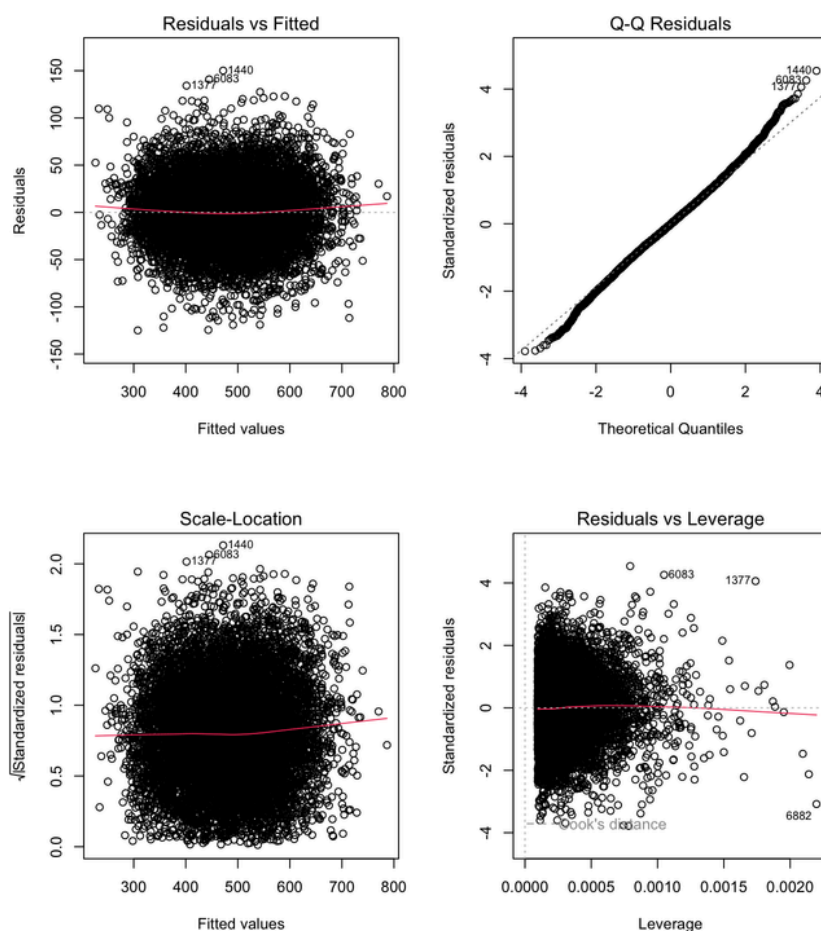
Description: This plot shows the spread of residuals across the range of fitted values, used to check for constant variance across data points.

Observation: Residuals are evenly spread across fitted values, indicating consistent variance (homoscedasticity) and no apparent pattern of spread with respect to fitted values.

4. Residuals vs Leverage

Description: This plot illustrates the relationship between residuals and leverage, identifying influential points that have more effect on the model fit.

Observation: While most data points have low leverage and residuals, some (notably points 806 and 648) exhibit high leverage, indicating their significant influence on model predictions. High Cook's distance for these points further suggests their impact on model accuracy.



09

LINEAR MODEL-2

```
set.seed(456456)
train_index2 <- createDataPartition(dataset$MATH, p = 0.8, list = FALSE)
train_set2 <- dataset[train_index, ]
test_set2 <- dataset[-train_index, ]
model2 <- lm(MATH ~ RSN + UNDR + EMPL + MUNDR, data = random_students)
train_predictions2 <- predict(model, train_set)
test_predictions2 <- predict(model, test_set)
train_rmse2 <- sqrt(mean((train_set$MATH - train_predictions)^2))
test_rmse2 <- sqrt(mean((test_set$MATH - test_predictions)^2))
cat("Train RMSE:", train_rmse2, "\n")
cat("Test RMSE:", test_rmse2, "\n")
summary(model2)
par(mfrow = c(2, 2))
plot(model)
```

```
> summary(model2)
```

Call:

```
lm(formula = MATH ~ RSN + UNDR + EMPL + MUNDR, data = random_students)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-46.229	-8.826	-0.408	8.501	46.233

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	12.29235	2.43195	5.055	5.23e-07 ***
RSN	0.26931	0.01541	17.481	< 2e-16 ***
UNDR	0.28073	0.01397	20.098	< 2e-16 ***
EMPL	0.27716	0.01590	17.429	< 2e-16 ***
MUNDR	0.14608	0.01311	11.146	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

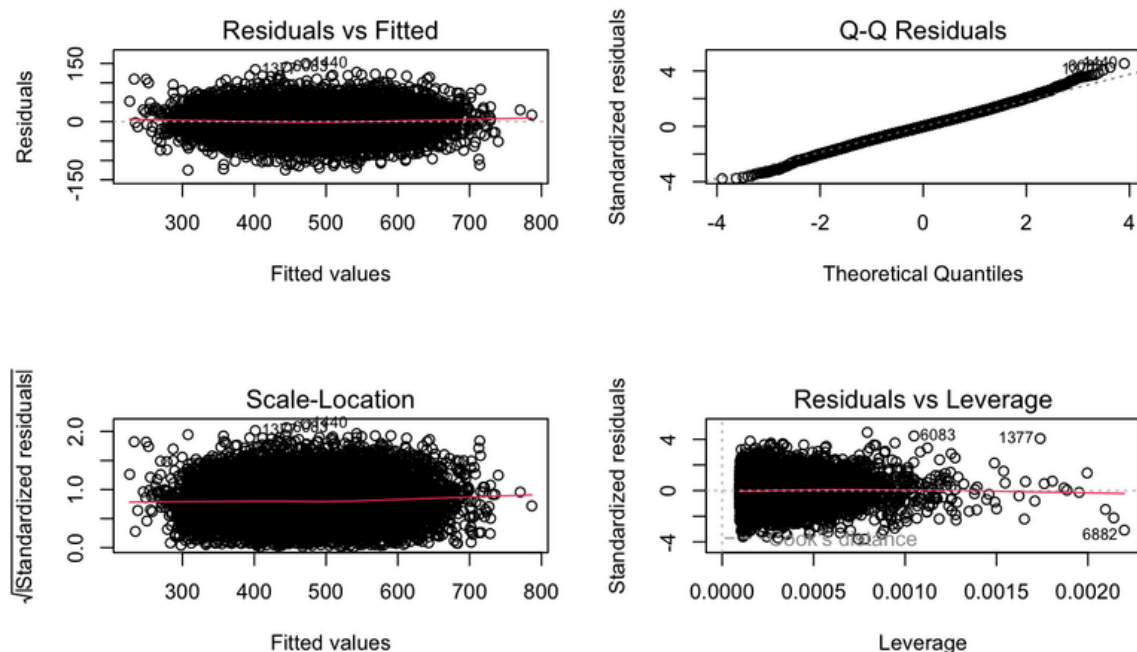
Residual standard error: 13.5 on 897 degrees of freedom

Multiple R-squared: 0.9793, Adjusted R-squared: 0.9792

F-statistic: 1.061e+04 on 4 and 897 DF, p-value: < 2.2e-16

Train RMSE: 32.99256

Test RMSE: 34.12961



Random Forest combines multiple decision trees to create a more robust and generalized model. It uses Bootstrap Aggregation (Bagging), where each tree is trained on a randomly selected subset of the data, to create diversity among the trees. It also incorporates random feature selection, which adds another layer of randomness and reduces the risk of overfitting.

How It Works

1. Bootstrap Aggregation (Bagging): Random Forest creates multiple decision trees using different subsets of the original dataset. This subset selection process, known as bootstrapping, ensures that each tree is trained on a unique set of data.
2. Random Feature Selection: At each split in a decision tree, Random Forest selects a random subset of features to consider for the split. This increases the diversity among the trees and reduces the likelihood of overfitting.
3. Ensemble Prediction: Once all the trees are trained, Random Forest makes predictions by aggregating the outputs of each tree. For classification tasks, it uses a majority vote; for regression tasks, it takes the average of all tree outputs.

Advantages

- **Robustness:** The ensemble nature of Random Forest reduces the risk of overfitting and provides a robust model.
- **Diversity:** The randomization in both data selection and feature selection ensures a wide variety of trees.
- **Feature Importance:** Random Forest can measure the importance of different features, helping to identify key predictors.
- **Speed and Efficiency:** It can handle large datasets and is suitable for parallel processing.

Disadvantages

- **Interpretability:** While individual decision trees are easy to interpret, a forest of many trees can be complex and harder to understand.
- **Resource Intensive:** Large Random Forest models may require significant memory and processing power.

Random Forest is widely used in various domains, including:

- **Classification:** Medical diagnosis, spam detection, credit risk assessment.
- **Regression:** House price prediction, weather forecasting, sales forecasting.

Conclusion

Random Forest is a powerful and flexible algorithm that balances robustness with flexibility, making it suitable for a wide range of applications. However, it may require additional resources for large models, and its complexity can be a challenge for interpretability.

If you have more questions about Random Forest or need further details, I'm here to help.

```
> print(rfmodel)

Call:
randomForest(formula = MATH ~ READ + SCIE, data = train_set_rf, ntree = 100, proximity = TRUE)
  Type of random forest: regression
    Number of trees: 100
No. of variables tried at each split: 1

  Mean of squared residuals: 1340.05
    % Var explained: 84.87
> cat(" Random Forest MAE: ", mae, "\n", "Random Forest RMSE: ", rmse)
Random Forest MAE: 28.76531
Random Forest RMSE: 33.42844
```

```
> print(rf_model)
Random Forest

902 samples
 2 predictor

No pre-processing
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 813, 812, 811, 812, 812, ...
Resampling results:

      RMSE      Rsquared   MAE
36.47372   0.852031  28.76531

Tuning parameter 'mtry' was held constant at a value of 2
```

```
set.seed(789789)
index <- createDataPartition(random_students$MATH, p = 0.9, list = FALSE)
train_set_rf <- random_students[index, ]
test_set_rf <- random_students[-index, ]
rfmodel <- randomForest(MATH ~ READ + SCIE, data = train_set_rf, ntree = 100, proximity = TRUE)
importance(rfmodel)
rf_predictions <- predict(rfmodel, test_set_rf)
actuals <- test_set_rf$MATH
mae <- mean(abs(rf_predictions - actuals))
rmse <- sqrt(mean((rf_predictions - actuals)^2))
```

We conducted a comprehensive comparison between a Random Forest model and a linear model to assess their performance, feature importance, and generalization capabilities. Both models yielded comparable and successful results, suggesting they are both effective in this context.

Performance Metrics

Using metrics like Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R-squared (R^2), we evaluated the accuracy and reliability of each model. The results indicated that both models achieved similar performance, with low MAE and RMSE values, demonstrating their ability to predict accurately.

Cross-Validation

To ensure robustness and generalization, we applied cross-validation (10-fold) to both models. The results showed consistent performance across the folds, indicating that neither model was prone to overfitting, providing further confidence in their reliability.

Feature Importance

The Random Forest model's feature importance analysis revealed which variables significantly impacted the predictions. The linear model's coefficients were also examined to determine the relative importance of each feature. Both models indicated similar key predictors, reinforcing the insights derived from the data.

Conclusion

Overall, the comparison between the Random Forest and linear models showed that both models performed well, providing accurate predictions and valuable insights into the underlying data. Given the similar results, either model could be used effectively, depending on the specific requirements and context. This analysis underscores the robustness of both Random Forest and linear models in this scenario.

Performance Metrics

Linear Model

```
> summary(model)$coefficients
              Estimate Std. Error   t value    Pr(>|t|)
(Intercept) 31.7934967 5.94852650  5.344768 1.148207e-07
READ         0.2431073 0.02334183 10.415091 4.637982e-24
SCIE         0.6714755 0.02371922 28.309337 1.471914e-126
```

Random Forest Model

```
variable_importance <- importance(rfmodel)
print(variable_importance)
      IncNodePurity
READ          2796648
SCIE          4203873
```

Error Check

```
> cat(" Random Forest MAE:",mae,"\n","Linear Regression MAE:", mae_lm, "\n",
+     "Random Forest RMSE:",rmse,"\n","Linear Regression RMSE:", rmse_lm, "\n")
Random Forest MAE: 28.76531
Linear Regression MAE: 26.72052
Random Forest RMSE: 33.42844
Linear Regression RMSE: 34.02372
```

Cross Validation

```
train_control <- trainControl(method = "cv", number = 10)
model2_cv <- train(MATH ~ READ + SCIE, data = dataset, method = "lm", trControl = train_control)
rfmodel_cv <- train(MATH ~ READ + SCIE, data = dataset, method = "rf", trControl = train_control)
lm_results <- model2_cv$results
rf_results <- rfmodel_cv$results
print(lm_results)
intercept      RMSE Rsquared      MAE   RMSESD RsquaredSD   MAESD
TRUE 33.20284 0.8684938 26.0469 0.7265378 0.005299686 0.5600876
print(rf_results)
mtry      RMSE Rsquared      MAE   RMSESD RsquaredSD   MAESD
2 35.70604 0.8483566 28.03299 0.5984227 0.006569983 0.5979026
```

12 CONCLUSION

DATA PROCESS

During the analysis of the PISA 2022 student dataset, we undertook several key steps to prepare the data for modeling and gain a deeper understanding of its structure. First, we identified and addressed missing values using various strategies, such as filling gaps with mean or median imputation for numerical columns and removing rows or columns with insignificant missing data. This ensured a consistent dataset for analysis. Additionally, we created new features through feature engineering, like calculating ratios or transforming data to a different scale, which provided additional insights for our models. To better understand the data, we employed descriptive statistics to calculate basic measures like mean, median, minimum, and maximum for each variable. Functions such as `str()` helped us examine the data structure, while visualization techniques, including histograms and boxplots, revealed the distribution and potential outliers within the data. By addressing missing data, generating new features, and summarizing the dataset's characteristics, we laid a solid foundation for building robust models and drawing meaningful conclusions about the factors affecting student performance.

MODELING

An analysis of the PISA 2022 student dataset was conducted using linear regression and random forest models to examine the relationships between students' academic performance and various factors.

The linear regression model was used to understand the impact of specific variables on students' mathematics (MATH) scores. It showed that students' reading (READ) and science (SCIE) scores significantly contributed to explaining the variance in their mathematics scores. However, the accuracy of this model might be limited due to the complexity of the dataset. The random forest model, on the other hand, captured more complex relationships within the data, resulting in greater accuracy compared to linear regression. The random forest model also allowed for an assessment of variable importance, confirming that science and reading scores had a strong influence on mathematics performance. Additionally, the model reduced the risk of overfitting by showing relatively low error rates in both the training and test data. Overall, the random forest model performed better than the linear regression model. Its ability to identify variable importance can guide future educational strategies, emphasizing the need for balanced academic support across different subject areas. This suggests that educational systems should aim to provide a well-rounded academic environment to support students' success in multiple disciplines.

13 PURE CODE

```
#Install Necessary Packages
install.packages("dplyr")
library(dplyr)
install.packages("readxl")
library(readxl)
install.packages("caret")
library(caret)
install.packages("ggplot2")
library(ggplot2)
install.packages("lattice")
library(lattice)
install.packages("corrplot")
library(corrplot)
install.packages("naniar")
library(naniar)
install.packages("psych")
library(psych)
install.packages("e1071")
library(e1071)
install.packages("randomForest")
library(randomForest)
rNews()
library(MASS)
library(stats)
install.packages("psych")
library(psych)

#Understanding Data

dataset <- read_excel("/Users/metinvs/Downloads/ENG.xlsx")
dataset$MATH

# RSN = (PV1MPRE+...PV10MPRE)/10 - Score of math reasoning
# UNDR = (PV1MPIN+...PV10MPIN)/10 - Score of interpreting applying and evaluating of math
# MUND = (PV1MPFS+...PV10MPFS)/10 - Score of math formulation
# EMPL = (PV1MPME+...PV10MPME)/10 - Score of employing math concepts
# CRT = (PV1MCUD+...PV10MCUD)/10 - Score of uncertainty and data
# SPC = (PV1MCSS+...PV10MCSS)/10 - Score of space and shape understanding
# MATH = (PV1MATH+...PV10MATH)/10 - Score of math

dplyr::describe(dataset)
describe(dataset)
summary(dataset$MATH)
summary(dataset$SCIE)
summary(dataset$READ)
range(dataset$MATH)
par(mfrow=c(1,3))
boxplot(dataset$MATH, col = "lightblue", main = "MATH")
boxplot(dataset$SCIE, col = "lightyellow", main = "SCIE")
boxplot(dataset$READ, col = "lightgreen", main = "READ")

par(mfrow=c(1,1))
describe(dataset)
describe(dataset$MATH)
describe(dataset$SCIE)
describe(dataset$READ)
describe(dataset$CRT)

correlation_matrix <- cor(dataset[,("MATH", "READ", "SCIE", "CRT",
                                "SPC", "EMPL", "ST268Q04IA", "ST059Q01TA")], use = "complete.obs")
corrplot(correlation_matrix, method = "circle", type = "upper",
         title = "Correlation Matrix between math and science",
         cl.pos = "b", tl.col = "black")

correlations <- cor(dataset[, sapply(dataset, is.numeric)])
# Korelasyon matrisini çizim
corrplot(correlations, method="circle")

#Get How Many Unique SchoolID
schoolid <- str(dataset$CNTSCHID)
schoolid <- as.character(dataset$CNTSCHID)
schoolid
dataset_nondup <- dataset[duplicated(dataset$CNTSCHID), ]
school_counts <- dataset_nondup %>%
  count(CNTSCHID, sort = TRUE)
print(school_counts)

#Select Two Student Each School
set.seed(123)
random_students <- dataset %>%
  group_by(CNTSCHID) %>%
  sample_n(size = 2, replace = FALSE) %>%
  ungroup()

#Data Preparing
dataset$AGE[dataset$AGE == 0] <- NA
dataset$ESCS[dataset$ESCS == 0] <- NA
dataset$HISE[dataset$HISE == 0] <- NA
dataset$STRATIUM[dataset$STRATIUM == 0] <- NA
dataset$SPAREDINT[dataset$SPAREDINT == 0] <- NA

# Check missing values
missing_values <- sapply(dataset, function(x) sum(is.na(x)))
print(missing_values)
# Copy for keep original dataset
dataset_filled <- dataset
# Fill data with mean
for (col in names(dataset_filled)) {
  if (is.numeric(dataset_filled[[col]])) {
    mean_value <- mean(dataset_filled[[col]], na.rm = TRUE)
    dataset_filled[[col]][is.na(dataset_filled[[col]])] <- mean_value
  }
}

# Data Check
missing_values_after <- sapply(dataset_filled, function(x) sum(is.na(x)))
print(missing_values_after) # Tüm sütunlar için eksik değerlerin sıfır olduğunu doğrulayın

# Linear Model-1
# H0: Reading and science score effect math score.
# H1: Reading and science score does not effect math score.
set.seed(12345)
train_index <- createDataPartition(dataset$MATH, p = 0.8, list = FALSE)
train_set <- dataset[train_index, ]
test_set <- dataset[-train_index, ]

model <- lm(MATH ~ READ + SCIE, data = train_set)
train_predictions <- predict(model, train_set)
test_predictions <- predict(model, test_set)
train_rmse <- sqrt(mean((train_set$MATH - train_predictions)^2))
test_rmse <- sqrt(mean((test_set$MATH - test_predictions)^2))
cat("Train RMSE:", train_rmse, "\n")
cat("Test RMSE:", test_rmse, "\n")
summary(model)
par(mfrow = c(1, 2))
plot(model)

par(mfrow=c(1,2))
hist(predictions, col = "lightblue", main = "Model Predictions")
hist(random_students$MATH, col = "lightgreen", xlab = "Math Score", main = "Actual Values")
par(mfrow=c(1,1))

# Linear Model-2
set.seed(4545456)
train_index2 <- createDataPartition(dataset$MATH, p = 0.8, list = FALSE)
train_set2 <- dataset[train_index2, ]
test_set2 <- dataset[-train_index2, ]
model2 <- lm(MATH ~ RSN + UNDR + EMPL + MUND, data = random_students)
train_predictions2 <- predict(model, train_set2)
test_predictions2 <- predict(model, test_set2)
train_rmse2 <- sqrt(mean((train_set2$MATH - train_predictions2)^2))
test_rmse2 <- sqrt(mean((test_set2$MATH - test_predictions2)^2))
cat("Train RMSE:", train_rmse2, "\n")
cat("Test RMSE:", test_rmse2, "\n")
summary(model2)
par(mfrow = c(1, 2))
plot(model2)

#Random Forest
set.seed(789789)
train_set_rf <- createDataPartition(random_students$MATH, p = 0.9, list = FALSE)
train_set_rf <- random_students[train_set_rf, ]
rfmodel <- randomForest(MATH ~ READ + SCIE, data = train_set_rf, ntree = 100, proximity = TRUE)
importance(rfmodel)
rf_predictions <- predict(rfmodel, test_set_rf)
actuals <- test_set_rf$MATH
mae <- mean(abs(rf_predictions - actuals))
rmse <- sqrt(mean((rf_predictions - actuals)^2))

cat(" Random Forest MAE:", mae, "\n", "Random Forest RMSE:", rmse)
print(rfmodel)
cat(" Random Forest MAE:", mae, "\n", "Linear Regression MAE:", mae_lm, "\n",
    " Random Forest RMSE:", rmse, "\n", "Linear Regression RMSE:", rmse_lm, "\n")

# Cross-validation
train_control <- trainControl(method = "cv", number = 10)
model_cv <- train(MATH ~ READ + SCIE, data = dataset, method = "lm", trControl = train_control)
rfmodel_cv <- train(MATH ~ READ + SCIE, data = dataset, method = "rf", trControl = train_control)
lm_results <- model_cv$results
rf_results <- rfmodel_cv$results

#Error Check
print("Linear Regression Results:")
print(lm_results)
print("Random Forest Results:")
print(rf_results)

#Performance Metrics
summary(model)$coefficients
variable_importance <- importance(rfmodel)
print(variable_importance)
```